



并行与分布式计算

Parallel & Distributed Computing


陈鹏飞
计算机学院



中山大学
SUN YAT-SEN UNIVERSITY

计算机学院 (软件学院)

SCHOOL OF COMPUTER SCIENCE AND ENGINEERING



Lecture 6 — Race Conditions and Synchronization in OpenMP

Pengfei Chen

School of Data and Computer Science



中山大學

SUN YAT-SEN UNIVERSITY

计算机学院 (软件学院)

SCHOOL OF COMPUTER SCIENCE AND ENGINEERING



Outline:

- **Correctness issues in parallel programming (in OpenMP)**
 - ❑ **Barriers (障碍, 屏障)**
 - ❑ **Examples of race conditions**
 - ❑ **Mutual Exclusion (互斥)**
 - ❑ **Memory fence (内存屏障)**



Concept: Synchronization

➤ **Synchronization**

- ❑ **The process of managing shared resources so that reads and writes occur in the correct order regardless of how the threads are scheduled**

➤ **Synchronization methods**

- ❑ **Barriers**

- ❑ **Mutual Exclusion (互斥) e.g. `pthread_mutex_lock`**

- ❑ **...**



Barriers in OpenMP

➤ Barrier

- ❑ A synchronization point at which every member in a team of threads must arrive before any member can proceed

➤ Syntax

#pragma omp barrier

- ❑ Automatically inserted at the end of worksharing constructs
- ❑ e.g., *for* pragma, *single* pragma, ...
- ❑ Can be disabled by using the *nowait* clause



Example: Use of Barrier

```
int numt = omp_get_num_threads();  
#pragma omp parallel shared(numt)  
{  
    int tid = omp_get_thread_num();  
    printf("hi, from %d\n", tid);  
    if (tid == 0) {  
        printf("%d threads say hi!\n", numt);  
    }  
}
```

Output using 4 threads

hi, from 3

hi, from 0

hi, from 2

4 threads say hi!

hi, from 1

◆ Question:

- What's the expected output?
- How can we let the last printf appear last?

Barbara Chapman, "A Guide to OpenMP," 2010.



Example: an Explicit Barrier

```
int numt = omp_get_num_threads();
#pragma omp parallel shared(numt)
{
    int tid = omp_get_thread_num();
    printf("hi, from %d\n", tid);
#pragma omp barrier
    if (tid == 0) {
        printf("%d threads say hi!\n", numt);
    }
}
```

Output using 4 threads

hi, from 3

hi, from 0

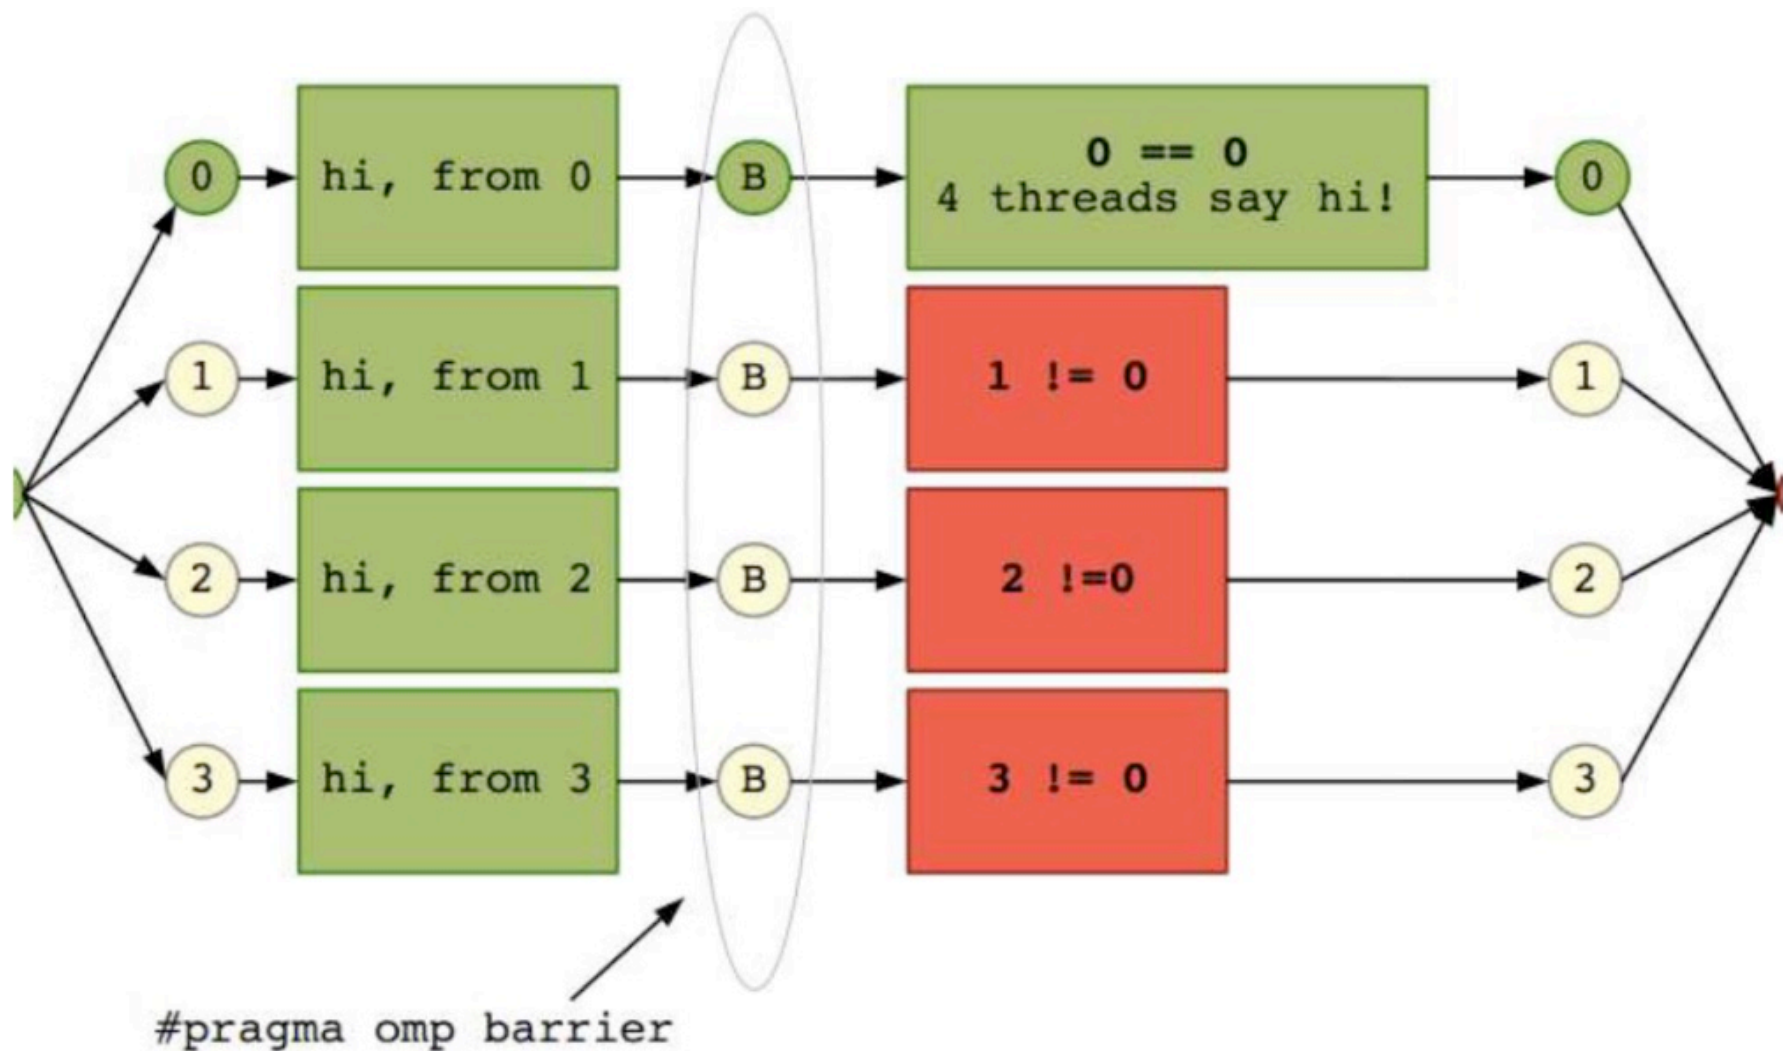
hi, from 2

hi, from 1

4 threads say hi!



Example: an Explicit Barrier





Clause: nowait

- The *nowait* clause tells the compiler that there is no need for a barrier

synchronization at the end of a *parallel for* loop or *single* block of code

Case: parallel, for, single Pragmas

```
for (i = 0; i < N; i++)  
    a[i] = alpha(i);  
if (delta < 0.0)  
    printf("delta < 0.0\n");  
for (i = 0; i < N; i++)  
    b[i] = beta(i, delta);
```



Solution: parallel, for, single Pragmas

```
#pragma omp parallel
{
    #pragma omp for nowait
    for (i = 0; i < N; i++)
        a[i] = alpha(i);
    #pragma omp single nowait
    if (delta < 0.0)
        printf("delta < 0.0\n");
    #pragma omp for
    for (i = 0; i < N; i++)
        b[i] = beta(i, delta);
}
```

← single 同一时间
只有一个线程执行

nowait 执行完不等



Mutual Exclusion

➤ **Mutual exclusion**

- ❑ **A kind of synchronization**

- ❑ **Allows only a single thread or process at a time to have access to shared resource**

- ❑ **Implemented using some form of locking**

➤ **Critical section (a high-level synchronization)**

- ❑ **Only one thread at a time will execute the structured block within a critical section**

➤ **Lock (a low-level synchronization)**



An Example of Race Condition

```
double area, pni, x;
int i, n;

...
area = 0.0;
for (i = 0; i < n; i++) {
    x = (i + 0.5) / n;
    area += 4.0 / (1.0 + x * x);
}
pi = area / n;
```

➤ What happens when we make the **for** loop parallel?




Race Condition

- A *race condition* is nondeterministic (非确定性) behavior caused by the times at which two or more threads access a shared variable
- For example, suppose both Thread A and Thread B are executing the statement...

```
area += 4.0 / (1.0 + x*x);
```




One Timing \Rightarrow Correct Sum





Another Timing \Rightarrow Incorrect Sum






Another Race Condition Example

```
struct Node {  
    struct Node* next;  
    int data;  
}  
  
struct List {  
    struct Node* head;  
};  
  
void AddHead (struct List* list, struct Node* node) {  
    node->next = list->head;  
    list->head = node;  
}
```





Thread 1 after Stmt. 1 of AddHead







Thread 2 Executes AddHead





Thread 1 after Stmt. 2 of AddHead





Why Race Conditions Are Nasty (令人讨厌)

- **Programs with race conditions exhibit nondeterministic behavior**
 - ❑ Sometimes give correct result
 - ❑ Sometimes give erroneous result
- **Programs often work correctly on trivial data sets and small number of threads**
- **Errors more likely to occur when number of threads and/or execution time increases**
- **Hence debugging race conditions can be difficult**



How to Avoid Race Conditions

- **Scope variables to be private to threads**
 - ▣ Use OpenMP *private* clause
 - ▣ Variables declared within threaded functions
 - ▣ Allocate on thread's stack (pass as parameter)
- **Control shared access with critical region**
 - ▣ Mutual exclusion and synchronization



Mutual Exclusion

- **We can prevent the race conditions described earlier ...**
 - ❑ **Ensure that only one thread at a time references or updates shared variables**
- **Mutual exclusion**
 - ❑ **A kind of synchronization**
 - ❑ **Allows only a single thread or process at a time to have access to shared resource**
 - ❑ **Implemented using some form of locking**



Do Flags Guarantee Mutual Exclusion?

```
int flag = 0;
```

```
void AddHead(struct List* list, struct Node* node) {  
    while (flag != 0) /* wait */;  
    flag = 1;  
    node->next = list->head;  
    list->head = node;  
    flag = 0;  
}
```



Flags Don't Guarantee Mutual Exclusion

```
int flag = 0;
```

flag

0

Thread 1

```
void AddHead(struct List* list, struct Node* node) {
```

```
    while (flag != 0) /* wait */;
```

```
    flag = 1;
```

```
    node->next = list->head;
```

```
    list->head = node;
```

```
    flag = 0;
```

```
}
```



Flags Don't Guarantee Mutual Exclusion

```
int flag = 0;
```

flag
0

Thread 1

Thread 2

```
void AddHead(struct List* list, struct Node* node) {
```

```
    while (flag != 0) /* wait */;
```

```
    flag = 1;
```

```
    node->next = list->head;
```

```
    list->head = node;
```

```
    flag = 0;
```

```
}
```




Flags Don't Guarantee Mutual Exclusion

```
int flag = 0;
```

flag
1

```
void AddHead(struct List* list, struct Node* node) {
```

```
    while (flag != 0) /* wait */;
```

```
    flag = 1;
```

```
    node->next = list->head;
```

```
    list->head = node;
```

```
    flag = 0;
```

```
}
```

Thread 1

Thread 2





Flags Don't Guarantee Mutual Exclusion

```
int flag = 0;
```

flag
1

```
void AddHead(struct List* list, struct Node* node) {  
    while (flag != 0) /* wait */;  
    flag = 1;  
    node->next = list->head;  
    list->head = node;  
    flag = 0;  
}
```

Thread 1

Thread 2





Flags Don't Guarantee Mutual Exclusion

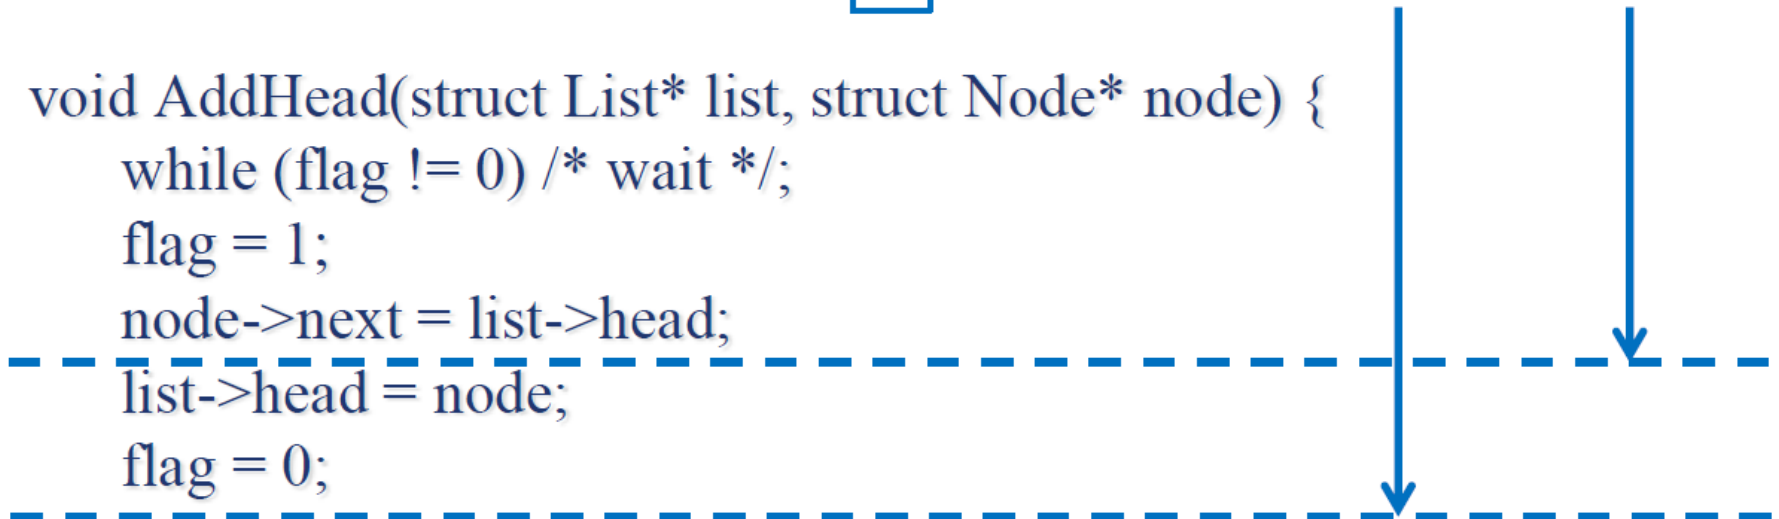
```
int flag = 0;
```

flag
0

```
void AddHead(struct List* list, struct Node* node) {  
    while (flag != 0) /* wait */;  
    flag = 1;  
    node->next = list->head;  
    list->head = node;  
    flag = 0;  
}
```

Thread 1

Thread 2





Flags Don't Guarantee Mutual Exclusion

```
int flag = 0;
```


flag
0

```
void AddHead(struct List* list, struct Node* node) {  
    while (flag != 0) /* wait */;  
    flag = 1;  
    node->next = list->head;  
    list->head = node;  
    flag = 0;  
}
```

Thread 1

Thread 2





Locking Mechanism

- The previous method fails because...
 - ❑ (i) Checking the value of *flag* and (ii) setting its value are two distinct operations
- We need some sort of *atomic* test-and-set
 - ❑ Operating systems provide functions to do this
- Lock
 - ❑ Synchronization mechanism used to control access to shared resources
 - ❑ (A generic term)



Pragma: critical

- **Critical section**
 - ❑ A portion of code that only one thread at a time may execute
(mutually exclusive)
- **Syntax in OpenMP**
 - `#pragma omp critical`
- **Good news! (^_^)**
 - ❑ Critical sections eliminate race conditions
- **Bad news! (@_@)**
 - ❑ Critical sections are executed sequentially
- **More bad news! (@_@)**
 - ❑ You have to identify critical sections yourself



Is the AddHead() Function Correct Now?

```
void AddHead(struct List* list, struct Node* node) {  
    node->next = list->head;  
    #pragma omp critical  
    list->head = node;  
}
```




中山大學

SUN YAT-SEN UNIVERSITY

计算机学院 (软件学院)

SCHOOL OF COMPUTER SCIENCE AND ENGINEERING



Is the AddHead() Function Correct Now?

- **You must protect both read and write access to any shared data**
- **For the AddHead() function, both lines need to be protected**



Corrected AddHead() Function

```
void AddHead(struct List* list, struct Node* node) {  
    #pragma omp critical  
    {  
        node->next = list->head;  
        list->head = node;  
    }  
}
```



OpenMP atomic (原子) Construct

- **Special case of a critical section to ensure atomic update to memory location**
- **Applies only to simple operations:**
 - ❑ **Pre- or post-increment (++)**
 - ❑ **Pre- or post-decrement (--)**
 - ❑ **Assignment with binary operator (of scalar types)**
- **Works on a single statement**

`#pragma omp atomic`

`counter += 5;`



Critical vs. Atomic

```
#pragma omp parallel for
{
    for (i = 0; i < n; i++) {
#pragma omp critical
        x[index[i]] += WorkOne(i);
        y[i] += WorkTwo(i);
    }
}
```

◆ Critical protect:

- Call to WorkOne()
- Finding value of index[i]
- Addition of x[index[i]] and results of WorkOne()
- Assignment to x array element

◆ Essentially, updates to elements in the x array are serialized

```
#pragma omp parallel for
{
    for (i = 0; i < n; i++) {
#pragma omp atomic
        x[index[i]] += WorkOne(i);
        y[i] += WorkTwo(i);
    }
}
```

◆ Atomic protects:

- Addition and assignment to x array element

◆ Non-conflicting updates with be done in parallel

◆ Protection needed only if there are two threads where the index[i] values match



```
#pragma omp parallel for
{
    for (i = 0; i < n; i++) {
#pragma omp critical
        x[index[i]] += WorkOne(i);
        y[i] += WorkTwo(i);
    }
}
```

A

```
#pragma omp parallel for
{
    for (i = 0; i < n; i++) {
#pragma omp atomic
        x[index[i]] += WorkOne(i);
        y[i] += WorkTwo(i);
    }
}
```

B



Critical vs. Atomic

```
#pragma omp parallel for
{
    for (i = 0; i < n; i++) {
#pragma omp critical
        x[index[i]] += WorkOne(i);
        y[i] += WorkTwo(i);
    }
}
```

◆ Critical protect:

- Call to WorkOne()
- Finding value of index[i]
- Addition of x[index[i]] and results of WorkOne()
- Assignment to x array element

◆ Essentially, updates to elements in the x array are serialized

```
#pragma omp parallel for
{
    for (i = 0; i < n; i++) {
#pragma omp atomic
        x[index[i]] += WorkOne(i);
        y[i] += WorkTwo(i);
    }
}
```

◆ Atomic protects:

- Addition and assignment to x array element

◆ Non-conflicting updates with be done in parallel

◆ Protection needed only if there are two threads where the index[i] values match



Summary

➤ Synchronization (in OpenMP)

□ Barrier

- Statement ordering among different threads
- Any statement after the barrier will be executed after the statements before the barrier in every thread

□ Mutual Exclusion

- Access ordering of shared resources
- A mechanism to avoid race conditions

□ Memory fence


- Data-related statement ordering in the same threads
- Any data-related statement before the memory fence will be executed before the statements after the memory fence



中山大學
SUN YAT-SEN UNIVERSITY

计算机学院 (软件学院)

SCHOOL OF COMPUTER SCIENCE AND ENGINEERING



Thank You !